

Python for HPC

Dr. Simone Bacchio

Computational Scientist

CaSToRC, The Cyprus Institute

National Competence Center in HPC



THE CYPRUS
INSTITUTE

RESEARCH • TECHNOLOGY • INNOVATION

Today's program

- 11:30-12:30 Performance in Python and Numpy
- 12:30-13:30 Lunch Break
- 13:30-14:30 Performance Optimization and Numba



Requirements:

- Some basic knowledge of Python
- Some basic knowledge of Numpy

Goal:

- Understand performance issues of Python and how to use it for HPC

What is your knowledge of Python??

Programming languages & Performance

- Not all programming languages are designed with performance in mind

Abstracted Programming Languages

- Python
- Matlab
- R, etc..

VS

HPC Programming Languages

- Fortran
- C, C++

Slow but easy-to-use




Fast but difficult

“**Pure**” Python is slow, very slow, but it can be made very fast... Very important to learn how!

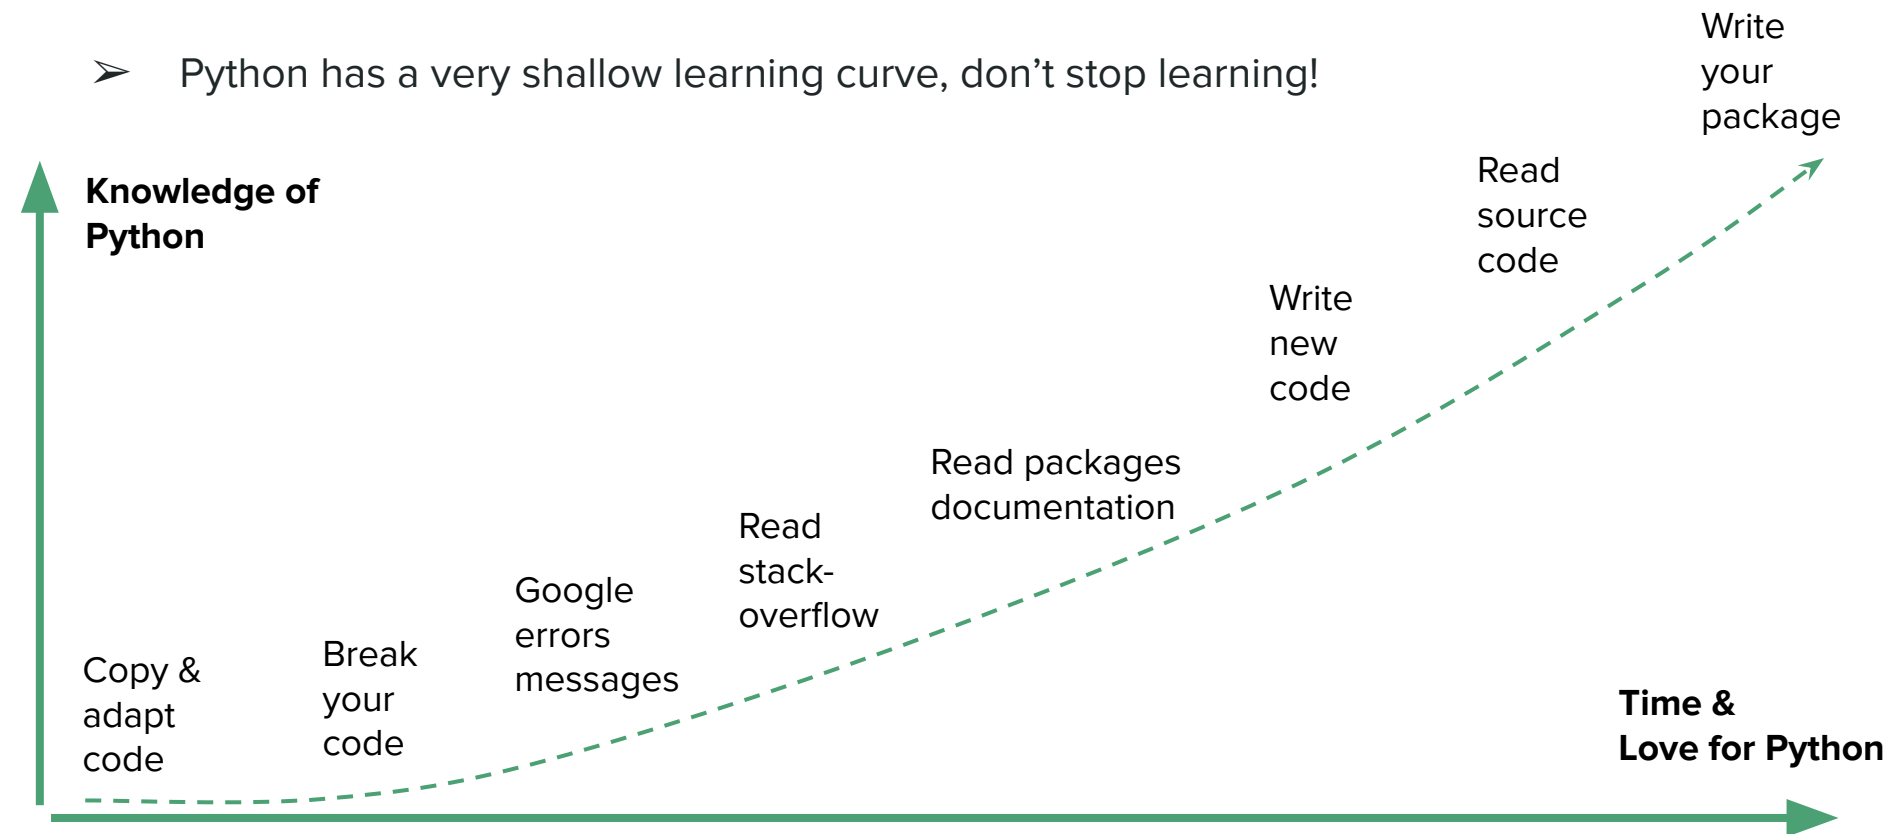
Why Python?

- **Most used programming language in data science**
- **Interpreted** and object oriented programming language
- Science- and data-oriented
- Easy to Learn and Use
- Huge community
- Hundreds of Python Libraries and Frameworks
- First choice for Big Data and **Machine learning**
- User-friendly and great **APIs**
- Easy deployment of software ([PyPI](https://pypi.org/))
- Build with a scientific approach ([PEPs](https://peps.python.org/))
- **Performance issues?** They can be overcome



How to learn Python?

- Python has a very shallow learning curve, don't stop learning!




How to use Python?



Pure Python and APIs

- Build up the logic and abstraction
- Make it effective and user-friendly
- Limit its use in computationally intensive parts



Compiled code & backends

- Many packages come with compile code
- Make it efficient and very fast (C performance)
- Use as much as possible in computations

Why is Python slow?

Python is a very powerful and flexible programming language, but...


- interpreted = bad (computational) performance
- it is important to know the strengths and the weaknesses!
- By its own it is not mean for High-Performance computing.

Built-in functions and HPC modules are based on **compiled** and **optimized** libraries.

Use as much as possible:

- built-in functions
- numerical modules ([Numpy](#), [Scipy](#), [Pandas](#), ...)
- compile your kernels ([Cython](#), [Pythran](#), [Numba](#), ...)

NEVER do for-loops on data!



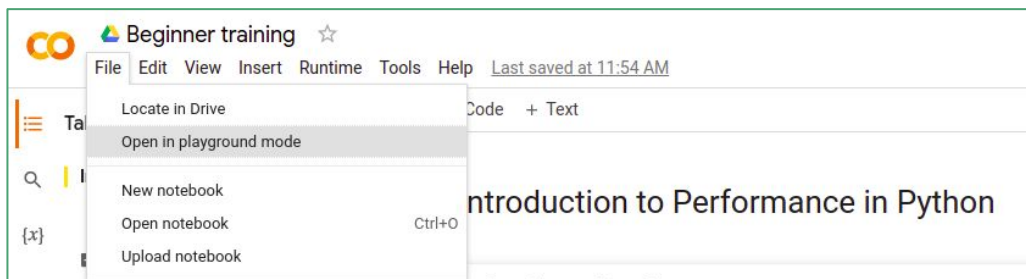
Numpy

- Numpy nowadays is the Python standard for numeric array calculations
- It is largely used and many packages are based on its **API**
 - **Scipy**: uses Numpy for implementing numerical algorithms
 - **Cupy**: a Numpy-compatible implementation for GPUs
 - **Numba**: JIT compiler for Python code using Numpy
 - **Pytorch**: its API is largely based on Numpy (not fully compatible tough)
 - ...
- A very good knowledge of Numpy is fundamental
- Documentation: <https://numpy.org/doc/stable/>
- Remaining of the training on Numpy



Let's get started

- For the training we will use Jupyter Notebooks in Google Colaboratory
https://colab.research.google.com/drive/1B9_gVPwIXohe2MqOJ5II_NI20sfQUldR
- Open the link and start a new notebook or open in playground mode

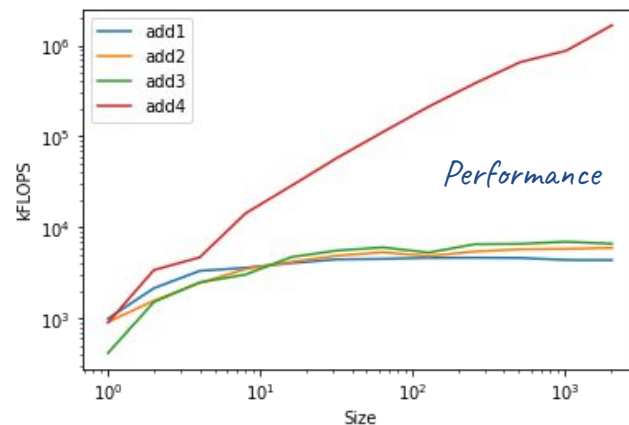



- Notebook and presentation also available on Github
<https://github.com/CaSToRC-Cyl/NCC-Beginner-Training-2022>

Performance



- For basic operations, Numpy achieves close-to-optimal performance and it is 1000x times faster than pure Python

Which one is Numpy??



Performance

- For basic operations, Numpy achieves close-to-optimal performance and it is 1000x times faster than pure Python




- **Remarks:**

- For small arrays Python overheads dominate
- Operations are done serially and between a step and another a new array is created

Introduction to Numpy

- The core of Numpy is `ndarray` (n-dimensional array)
- An `ndarray` is defined by
 - **shape**: the size of the array along each dimension
 - **dtype**: the data type of the array and its size (`arr.dtype.itemsize`)
 - **ordering**: the data ordering in memory (C or F-contiguous)
- Any operation on the array is done via compiled code with high performance
- Implementation-wise `ndarray` is a **view** of a 1-dimensional array (unrolled data)
 - See **Python Buffer Protocol**, <https://docs.python.org/3/c-api/buffer.html>
 - See **Array Interface Protocol**, <https://numpy.org/doc/stable/reference/arrays.interface.html>
 - See e.g. `arr.__array_interface__`

How does it work?



- N-dimensional arrays are *views* of **unrolled data**
- The shape is an artifact on the Python side but implementation-wise numpy always process unrolled data
- **NOTE:** for performance purposes, often many operation return different view of the same pointer. Therefore be careful when modifying arrays in-place!

Item access, modification and slicing

- Arrays elements can be accessed and modified as for lists
 - Elements per dimensions can be either extracted serially or at once
 - E.g. `arr[0,1,2,3] = arr[0][1][2][3]`
 - The first, of course, is optimal because avoids creation of intermediate arrays
- Slices, ranges or lists can be used for accessing multiple elements at once
 - Slices are open ranges
 - E.g. `:10 == 0:10`
 - **Note:** tuples cannot be used!
- Dimensions can be skipped using ellipses (`. . .`)
- Broadcasting also applies for element assignment
- Assignment and assigning operations (`+=`) might change the original array!

Universal functions

- See <https://numpy.org/doc/stable/reference/ufuncs.html>
- **Element-wise operations**
 - **Binary operations:** `+(add)`, `-(sub)`, `*(mul)`, `/(div)`, `%(mod)`, `==(eq)`, `** (pow)`, ...
 - **Math functions:** `exp`, `log`, `sin`, `cos`, `tan`, ...
 - Custom functions can be implemented via `np.vectorize`
- **Reductions**
 - Equal to: `for i in range(len(A)): r = op(r, A[i])`
 - Examples: `sum`, `mean`, `std`, `max`, `min`
 - They can be performed axis-wise (via argument `axis`)
 - Custom reductions can be implemented via `ufunc.reduce`
 - E.g. `sum = add.reduce`

Performance limitations

```
y = x ** 2 + 2 * x + 1
```

VS

```
for(int i=0; i<N; i++) {  
    y[i] = x[i] ** 2 + 2 * x[i] + 1  
}
```

What is the difference?

Performance limitations

$$\begin{array}{c} y = x ** 2 + 2 * x + 1 \\ \underbrace{\quad\quad\quad}_{a1} \quad + \quad \underbrace{\quad\quad\quad}_{a2} \\ \underbrace{\quad\quad\quad}_{a3} \quad + \quad 1 \\ \underbrace{\quad\quad\quad\quad\quad}_{y} \end{array}$$

VS

```
for(int i=0; i<N; i++) {  
    y[i] = x[i] ** 2 + 2 * x[i] + 1  
}
```

Left: 4 loop over data, 5 array access, 3 extra arrays allocated

Right: 1 loop over data, 1 array access, no extra array allocated


- Any operation on the arrays creates intermediate results and therefore new arrays
- This is quite a performance drawback because many allocations and loops are done
- Additionally a compiled loop can be optimized and use “special” operations
- This issue can be solved using `np.vectorize`

Broadcasting

- Arrays of different dimensions can be operated together

Requirements:

- Sizes must be either 1 or equal comparing from right to left
- **If same size:**
they are combined element-wise
- **If one-sized:**
same value used for all axis
- **If missing dimensions:**
automatically one-sized from left



Additional operations

- With numpy you can almost do everything, without having to write a for-loop in Python
- For this you need a good knowledge of the API and can be achieved only practicing!
- E.g. how to do “ $x[i+1] - x[i]$ ”? `y = np.roll(x, -1) - x`
 - See e.g. <https://numpy.org/doc/stable/reference/routines.array-manipulation.html>
- Many examples available online or on stack overflow... just search!

You didn't find what you are looking for?

- **Try Numba!**

Numba: a JIT compiler for Python

Numba is an open source JIT compiler that translates a subset of Python and NumPy code into fast machine code.

- Documentation: <https://numba.pydata.org>
- Installation: `pip install numba`
- CPU compiler: `from numba import jit`
- GPU API: `from numba import cuda`

Easy compilation and parallelization

- Numba easily compiles, vectorize and parallelize Python code!

Advantages?

- The code gets compiled reaching C-performance
- The code can run in parallel using multi-threading

Issue?

- You need to explicitly write for-loops in Python!

So if you do not have any other way than writing explicitly a for-loop...

Then do it and use Numba to speed it up!

```
from numba import njit, prange

@njit(parallel=True)
def difference(arr):
    N = arr.shape[0]
    out = np.empty_like(arr)
    for i in prange(N):
        out[i] = arr[(i+1)%N] - arr[i]
    return out
```

Conclusions

- Never do for-loop on data in Python
- Numpy comes first at rescue with its very user-friendly API
 - **NOTE:** Other packages are available, e.g. Pandas dataframe (on Wednesday) but a very good knowledge of numpy is fundamental
- Use Numba to speed-up Python code
 - We just had time to scratch the surface. Give it a try it is very useful!
 - More will be covered in the intermediate training including GPU programming

Questions??